

Project History

This project was created out of a codebase for Tarah Wheeler's doctoral dissertation. After a discussion between Jon Martin and Tarah Wheeler, we collaborated to build a project which could begin to detect relationships between text using well-tested and already available natural language processing tools. Though the logic and pseudocode began with Tarah's concepts for literary influence, the by-far-largest part of the engineering, optimization, and code development was done by Jon. At the point at which we chose to move the project away from Tarah's embedded dissertation data (some of which is private or not-for-public release, limiting the usefulness of the codebase to the public) into an MIT-licensed project intended for public release, Jon had the lion's share of commits. That history of the project should not be lost because it was done in a previous iteration.

May 8, 2022 – Jul 9, 2023

Contributions: Commits ▾

Contributions to main, excluding merge commits and bot accounts



As the project moves forward and (hopefully) other contributors can join in, please do send gratitude and respect to Jon's leading and massive contributions during the messiest creative part of the project.

Tarah Wheeler
July 9th, 2023